

Algorithm design: stable matching

Reading: Kleinberg & Tardos Ch. 1
rephrased version on moodle

Algorithm design process

Algorithm:

computational procedure correctly solving a problem

1. pose the problem rigorously
(formulate a *mathematically clean* definition)
2. propose a solution as a procedure
(description can be *implemented* within specified computational model, e.g., a random-access machine)
3. prove that the procedure correctly solves the problem
4. analyze the procedure's running **time**

Algorithm design process

Algorithm:

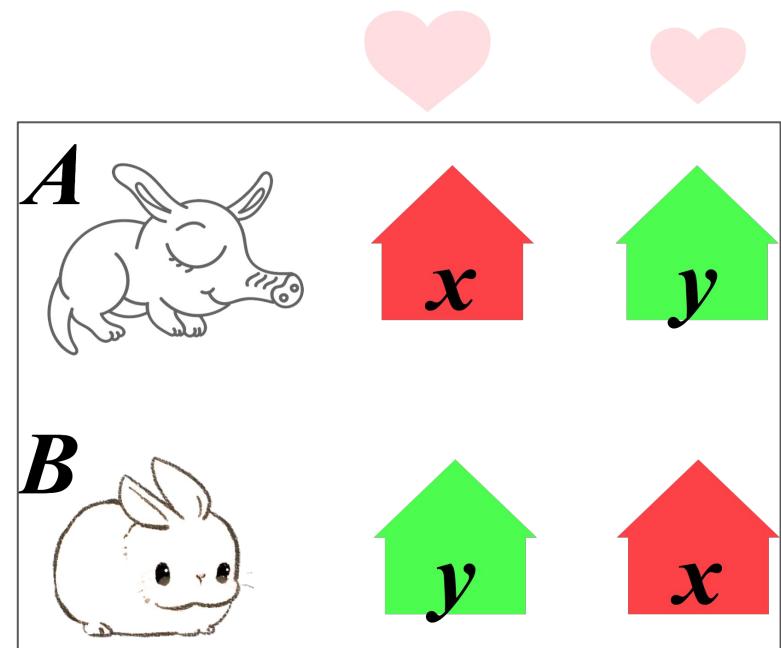
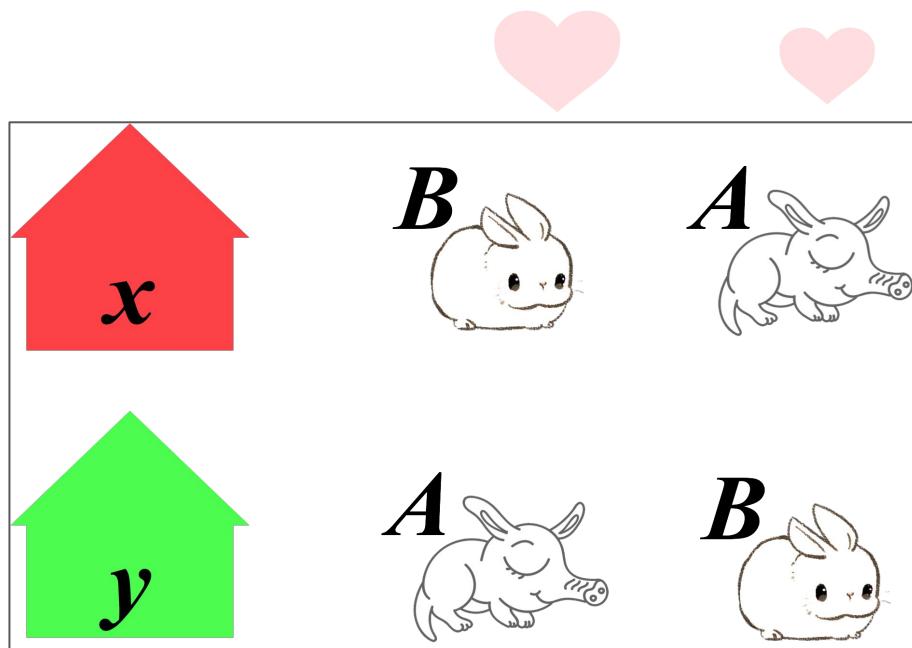
computational procedure correctly solving a problem

0. work examples to gain
familiarity/intuition

1. pose the problem rigorously
(formulate a *mathematically clean* definition)
2. propose a solution as a procedure
(description can be *implemented* within specified computational model, e.g., a random-access machine)
3. prove that the procedure correctly solves the problem
4. analyze the procedure's running **time**

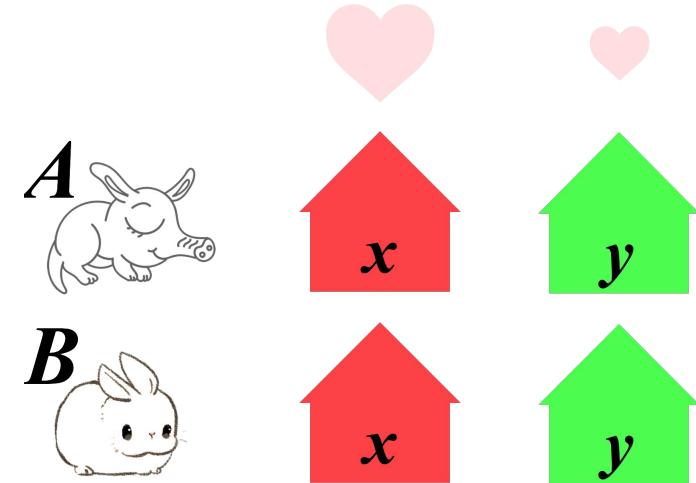
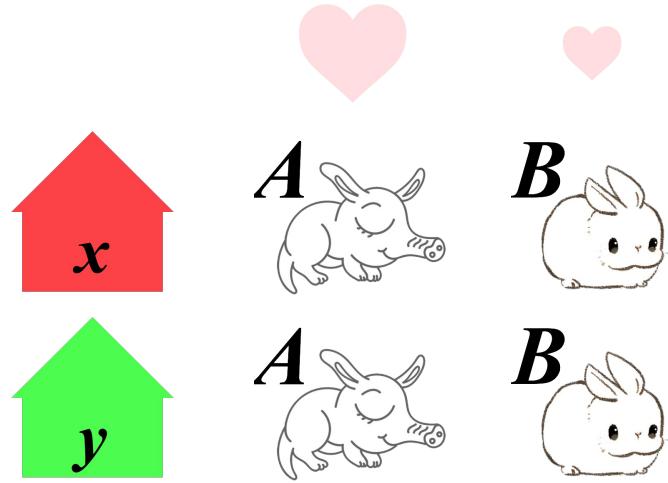
Problem setup: pet adoption

A shelter is looking to have its animals adopted into loving families. How do you decide which animals should become pets to which families?

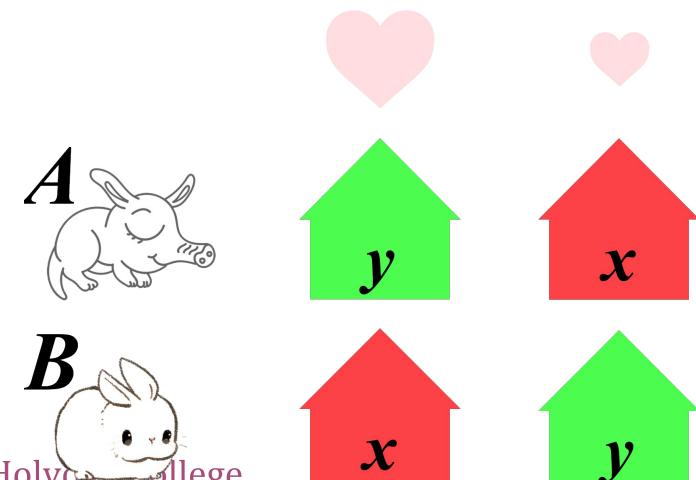
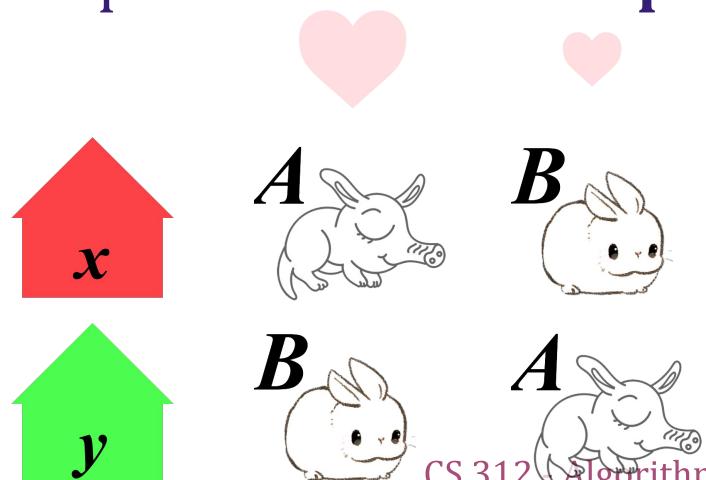


What adoptions would you suggest?

Example 1. Universal prefs



Example 2. Inconsistent prefs



Algorithm design process

Algorithm:

computational procedure correctly solving a problem

1. pose the problem rigorously
(formulate a *mathematically clean* definition)
2. propose a solution as a procedure
(description can be *implemented* within specified computational model, e.g., a random-access machine)
3. prove that the procedure correctly solves the problem
4. analyze the procedure's running **time**

Algorithm design process: step 1

Algorithm:

computational procedure correctly solving a problem

1. pose the problem rigorously
(formulate a *mathematically clean* definition)

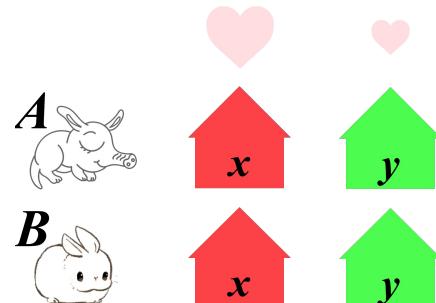
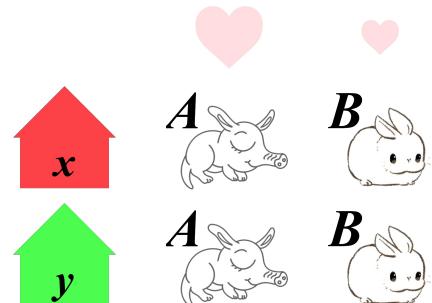
Formulating the problem of pet adoption

Given: n families looking to adopt n animals

- Each family ranks all animals
- Each animal ranks all families
- For simplicity, **assume each family can only adopt one animal**

Can we match animals to families s.t. everyone is “**happy**”?

- Not necessarily -- what if the aardvark were everyone’s top choice?



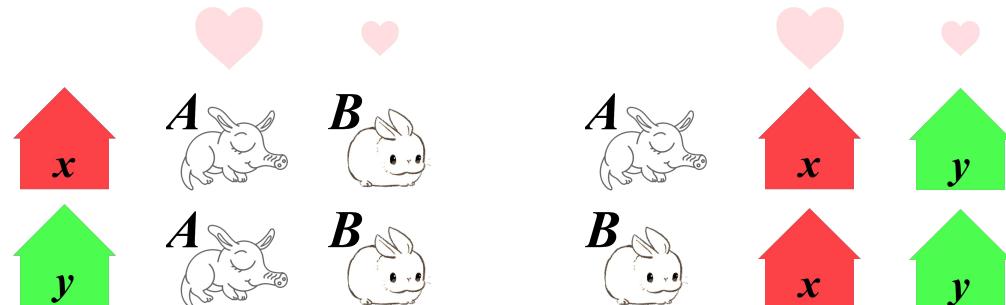
Stable matchings for pet adoption

Given: n families looking to adopt n animals

- Each family ranks all animals
- Each animal ranks all families
- For simplicity, assume each family can only adopt one animal

Can we match animals to families in a **stable** way?

- *Stable*: Don't match (f, j') and (f', j) if f and j would both prefer to be matched with each other



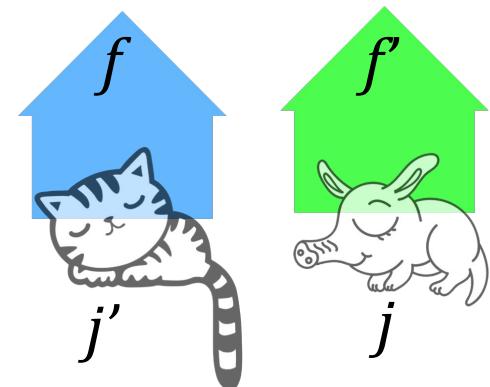
Step 1: precise problem formulation

Stable matching problem

Input: preference lists for n families and n animals

Output: ?? need definitions first

- **Matching:** set M of family-animal pairs
each family/animal participates in *at most* one pair
- **Perfect matching:** each family/animal in *exactly* one pair
- **Instability** or unstable pair (with respect to matching M):
a pair $(f, j) \notin M$ such that
 - $(f, j') \in M$ but f prefers j to j'
 - $(f', j) \in M$ but j prefers f to f'



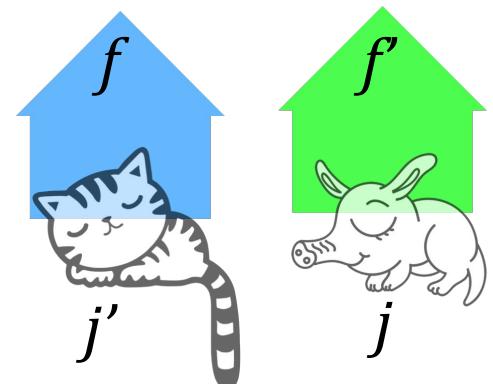
Step 1: precise problem formulation

Stable matching problem

Input: preference lists for n families and n animals

Output: ?? need definitions first

- **Matching:** set M of family-animal pairs
each family/animal participates in *at most* one pair
- **Perfect matching:** each family/animal in *exactly* one pair
- **Instability** or unstable pair (with respect to matching M):
a pair $(f, j) \notin M$ such that
 - $(f, j') \in M$ but f prefers j to j'
 - $(f', j) \in M$ but j prefers f to f'
- **Stable matching:**
perfect matching with no instabilities



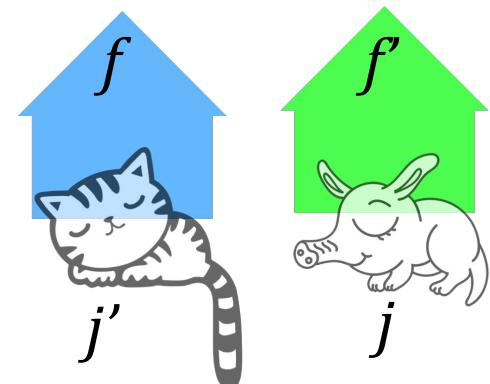
Step 1: precise problem formulation

Stable matching problem

Input: preference lists for n families and n animals

Output: a stable matching

- **Matching:** set M of family-animal pairs
 - each family/animal participates in *at most* one pair
- **Perfect matching:** each family/animal in *exactly* one pair
- **Instability** or unstable pair (with respect to matching M):
 - a pair $(f, j) \notin M$ such that
 - $(f, j') \in M$ but f prefers j to j'
 - $(f', j) \in M$ but j prefers f to f'
- **Stable matching:**
 - perfect matching with no instabilities

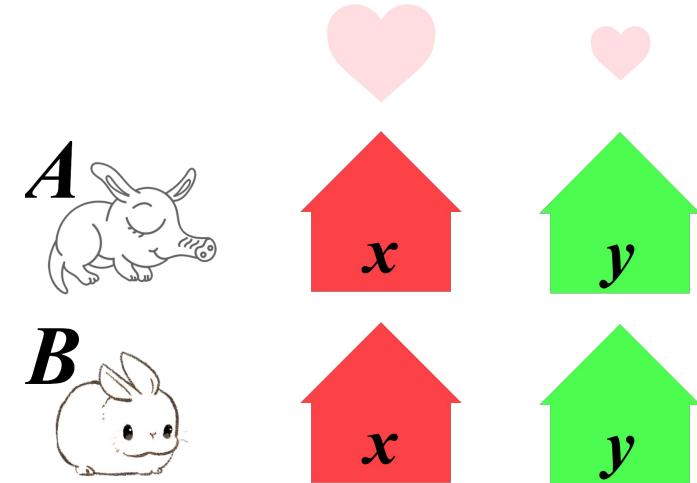
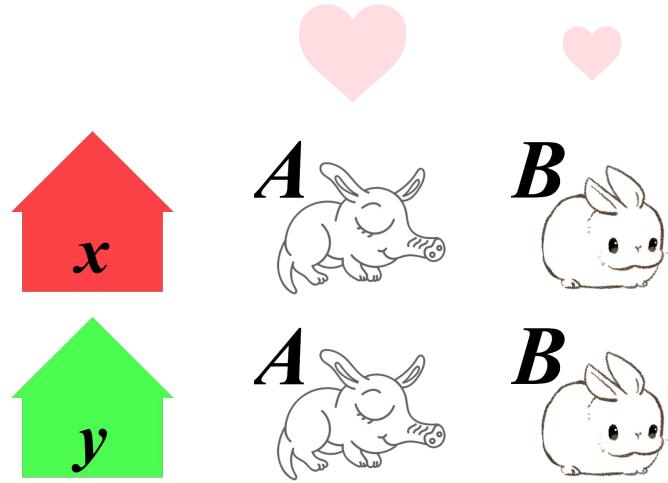


Stable matchings

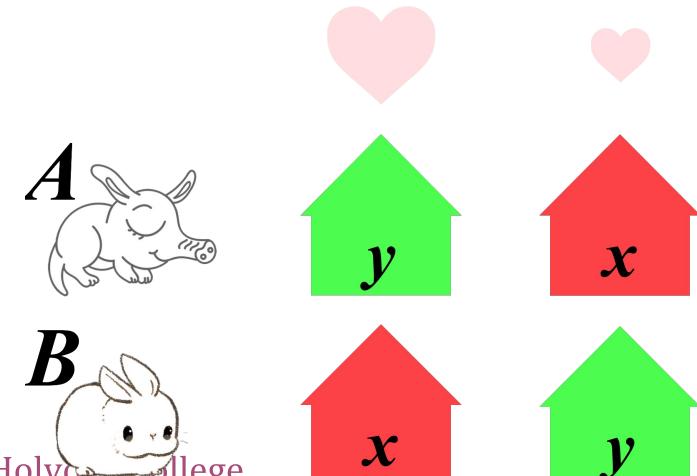
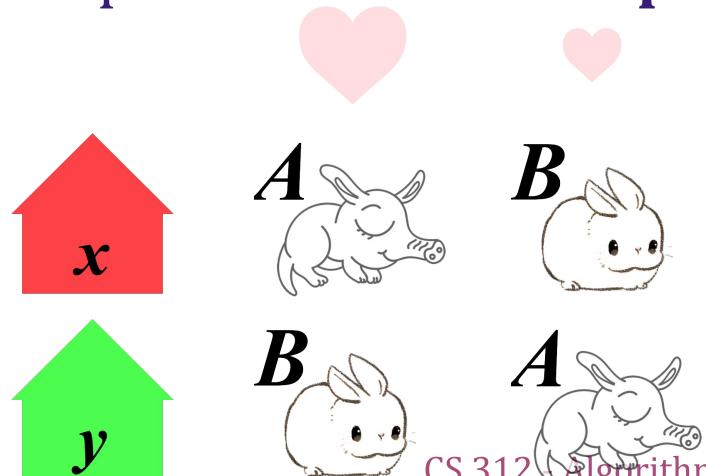
- Do they always exist?
- Are they unique?

What are the possible perfect matchings? Which are stable?

Example 1. Universal prefs



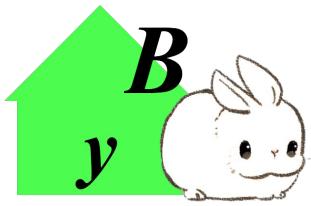
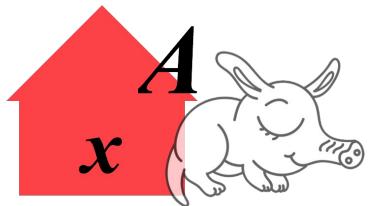
Example 2. Inconsistent prefs



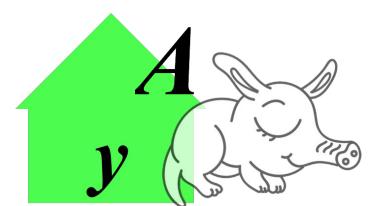
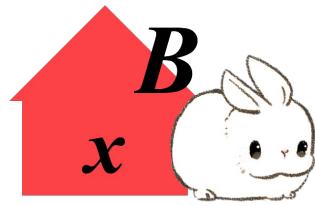
Perfect matchings

2 families, 2 animals

$$M = \{(x,A), (y,B)\}$$



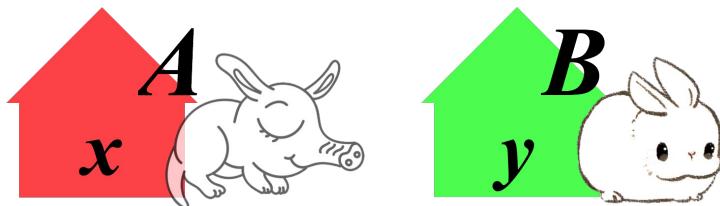
$$M' = \{(x,B), (y,A)\}$$



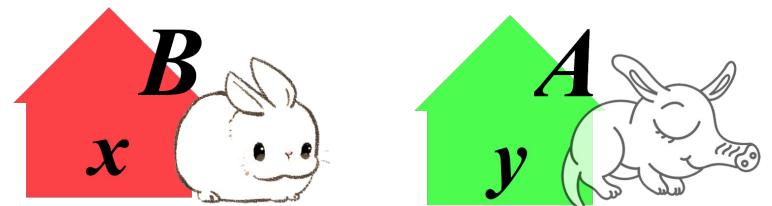
Stable?

2 families, 2 animals

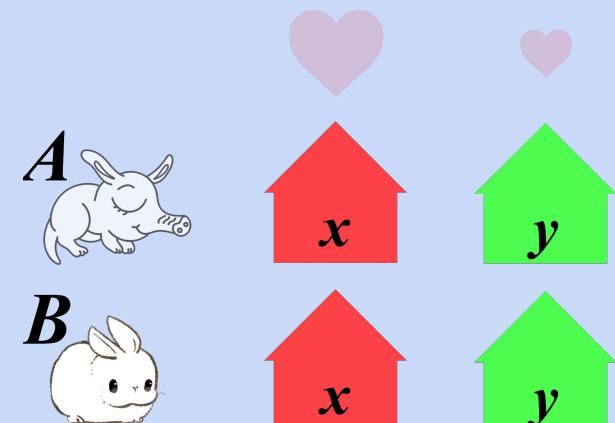
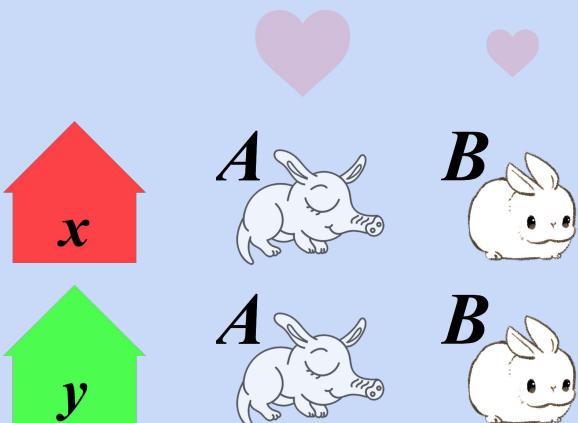
$$M = \{(x, A), (y, B)\}$$



$$M' = \{(x, B), (y, A)\}$$



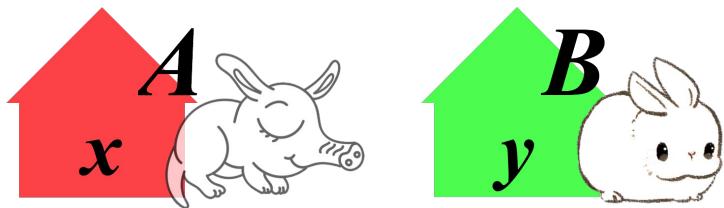
Universal preferences



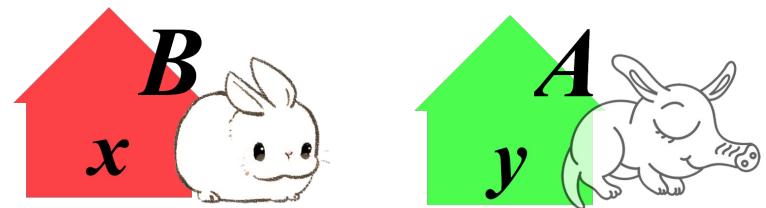
Stable?

2 families, 2 animals

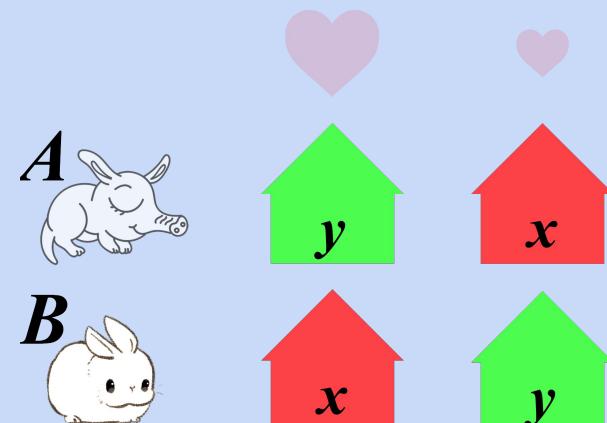
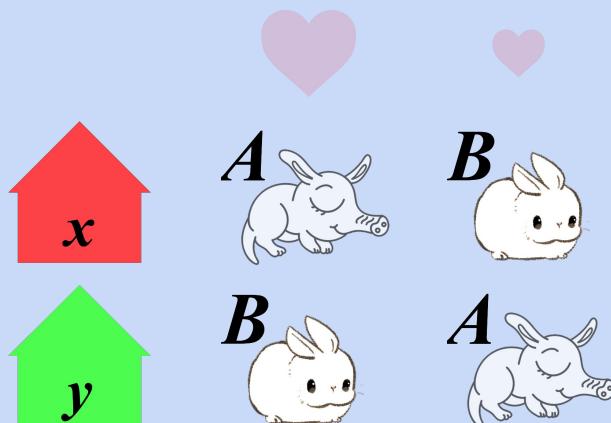
$$M = \{(x, A), (y, B)\}$$



$$M' = \{(x, B), (y, A)\}$$



Inconsistent preferences



Algorithm design process: step 2

Algorithm:

computational procedure correctly solving a problem

1. pose the problem rigorously
(formulate a *mathematically clean* definition)
2. propose a solution as a procedure
(description can be *implemented* within specified computational model, e.g., a random-access machine)

Algorithm design process: step 2

Algorithm:

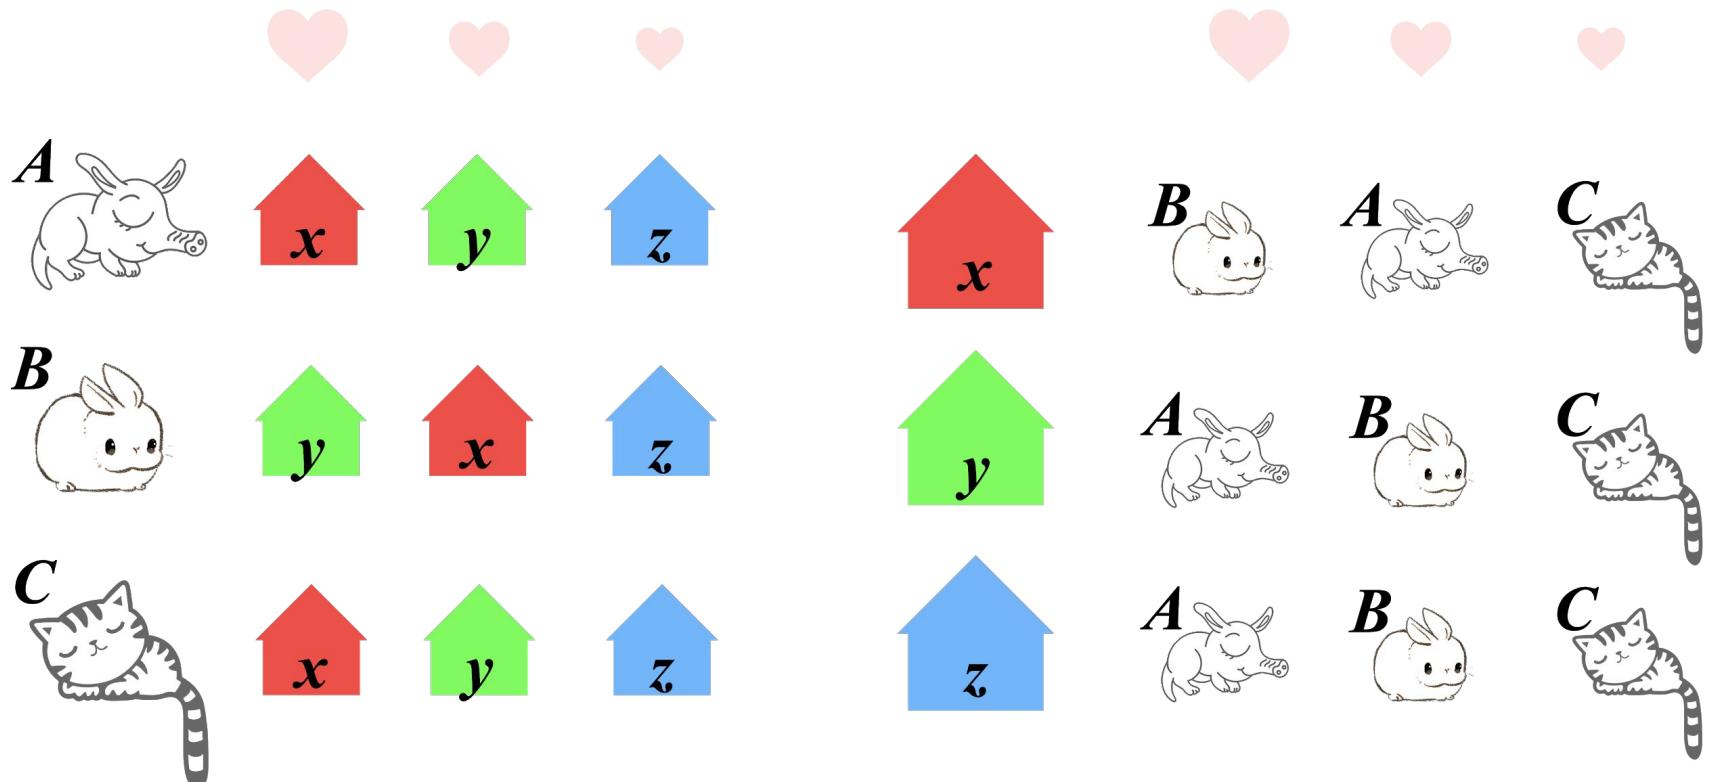
computational procedure correctly solving a problem

1. pose the problem rigorously
(formulate a *mathematically clean definition*)
2. propose a solution as a procedure
(description can be *implemented* within specified computational model, e.g., a random-access machine)

What would a “brute force” approach be?

Towards an algorithm: build M incrementally

A family walks into the shelter...



Propose-and-Reject (Gale-Shapley) Algorithm

Initialize all families and animals as **free**

while some family is **free** and hasn't proposed to every animal

 Choose such a family f

 Let a be the highest ranked animal to whom f has not yet proposed

 if a is **free**

f and a become **committed**

 else // a is **committed** to f'

 if a prefers f

f' becomes **free**

f and a become **committed**

 else

a rejects f and f remains **free**

Algorithm Design Process

Algorithm design process: analysis (steps 3 and 4)

Algorithm:

computational procedure correctly solving a problem

1. pose the problem rigorously
(formulate a *mathematically clean* definition)
2. propose a solution as a procedure
(description can be *implemented* within specified computational model, e.g., a random-access machine)
3. prove that the procedure correctly solves the problem
4. analyze the procedure's running **time**

Algorithm analysis: towards a proof

Some natural questions:

- Can we guarantee every family and animal gets a match?
- Can we guarantee the resulting adoptions are stable?
- Can we guarantee the algorithm terminates?

Some initial observations:

- (F1) Once matched, an animal stays matched and only "upgrades" during the algorithm.
- (F2) Each family proposes to animals in order of family's preferences.

Algorithm design process: analysis (steps 3 and 4)

Algorithm:

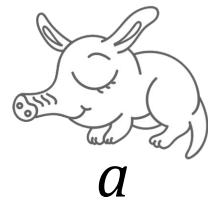
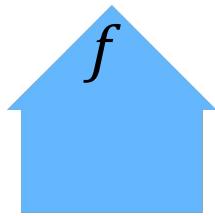
computational procedure correctly solving a problem

1. pose the problem rigorously
(formulate a *mathematically clean* definition)
2. propose a solution as a procedure
(description can be *implemented* within specified computational model, e.g., a random-access machine)
3. prove that the procedure correctly solves the problem
4. analyze the procedure's running **time**

Can we guarantee a perfect matching?

Yes! Proof by contradiction...

- Suppose not all families and animals have matches
 - Then there exists unmatched family f and unmatched animal a
- a was never matched during the algorithm (by F1)
- But f proposed to every animal (by termination condition)
- When f proposed to a , it was unmatched and yet rejected f
- Contradiction!

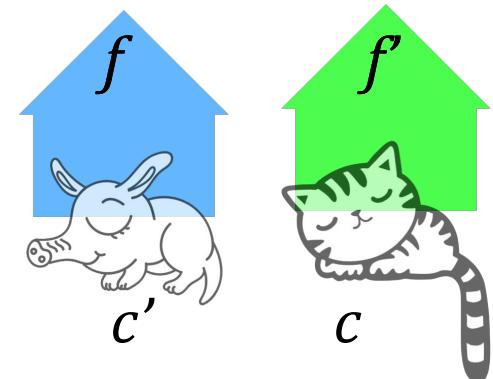


(F1) Once matched, an animal stays
matched and only “upgrades”

Can we guarantee the resulting adoptions are stable?

Yes! Proof by contradiction.

- Suppose there is an instability (f, c)
 - f is matched to c' but prefers c to c'
 - c is matched to f' but prefers f to f'

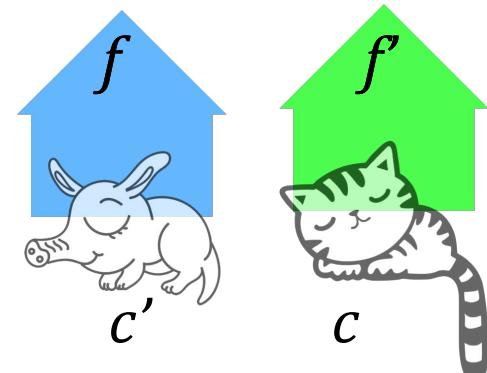


Can we guarantee the resulting adoptions are stable?

Yes! Proof by contradiction.

- Suppose there is an instability (f, c)
 - f is matched to c' but prefers c to c'
 - c is matched to f' but prefers f to f'
- By (F2), f must have proposed to c before proposing and becoming committed to c'

(F2) Each family proposes to animals in order of family's preferences.

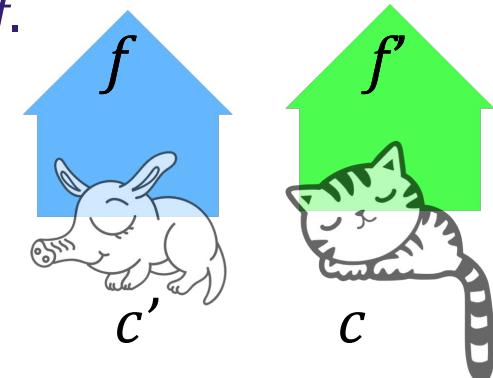


Can we guarantee the resulting adoptions are stable?

Yes! Proof by contradiction.

- Suppose there is an instability (f, c)
 - f is matched to c' but prefers c to c'
 - c is matched to f' but prefers f to f'
- By (F2), f must have proposed to c before proposing and becoming committed to c'
- Since f isn't matched to c at the end of the algorithm, it must have rejected f (either immediately or upon receiving a better proposal). By (F1), it prefers its final match f' to f .

(F1) Once matched, an animal stays
matched and only "upgrades"



Propose-and-Reject (Gale-Shapley) Algorithm

Initialize all families and animals as **free**

while some family is **free** and hasn't proposed to every animal

 Choose such a family f

 Let a be the highest ranked animal to whom f has not yet proposed

 if a is **free**

f and a become **committed**

 else // a is **committed** to f'

 if a prefers f

f' becomes **free**

f and a become **committed**

 else

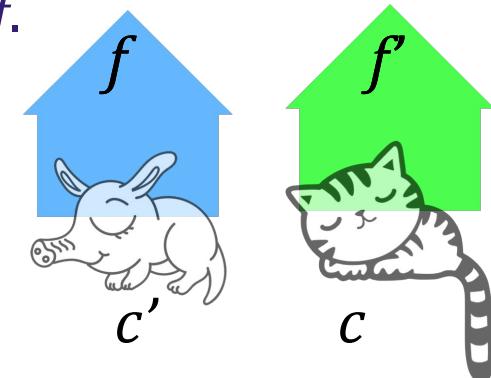
a rejects f and f remains **free**

Can we guarantee the resulting adoptions are stable?

Yes! Proof by contradiction.

- Suppose there is an instability (f, c)
 - f is matched to c' but prefers c to c'
 - c is matched to f' but prefers f to f'
- By (F2), f must have proposed to c before proposing and becoming committed to c'
- Since f isn't matched to c at the end of the algorithm, it must have rejected f (either immediately or upon receiving a better proposal). By (F1), it prefers its final match f' to f .

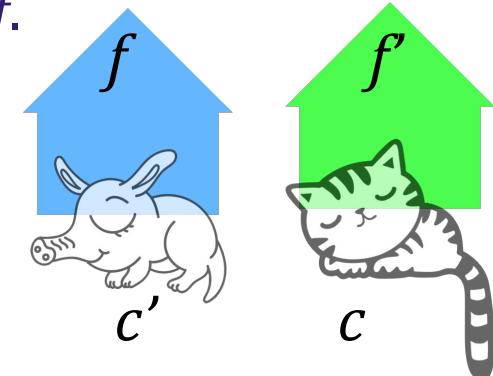
(F1) Once matched, an animal stays
matched and only "upgrades"



Can we guarantee the resulting adoptions are stable?

Yes! Proof by contradiction.

- Suppose there is an instability (f, c)
 - f is matched to c' but prefers c to c'
 - c is matched to f' but prefers f to f'
- By (F2), f must have proposed to c before proposing and becoming committed to c'
- Since f isn't matched to c at the end of the algorithm, it must have rejected f (either immediately or upon receiving a better proposal). By (F1), it prefers its final match f' to f .
- Contradiction!



Algorithm design process: analysis (steps 3 and 4)

Algorithm:

computational procedure correctly solving a problem

1. pose the problem rigorously
(formulate a *mathematically clean* definition)
2. propose a solution as a procedure
(description can be *implemented* within specified computational model, e.g., a random-access machine)
3. prove that the procedure correctly solves the problem
4. analyze the procedure's running **time**

Can we guarantee the algorithm terminates?

Yes! Proof...

- In every round, some family proposes to some animal that they haven't already proposed to
- n families and n animals \Rightarrow at most n^2 proposals
 \Rightarrow at most n^2 rounds of the algorithm

But what about running time?

Yes! Proof...

- In every round, some family proposes to some animal that they haven't already proposed to
- n families and n animals \Rightarrow at most n^2 proposals

\Rightarrow **at most n^2 rounds of the algorithm**

How do we formalize this in terms of **running time**?